


# Introduction to Containers and Kubernetes

Ruben Orduz + Nolan Brubaker  
T: @rdodev + @palendae


# A bit about Ruben

- Software developer
  - Have been using cloud tech in production for about 7 years
  - OpenStack developer advocate for a few years
  - Presently work as Field Engineer at Heptio
- 

# A bit about Nolan

- Software developer
- Background in Python web development
- OpenStack-Ansible contributor for 3 years
- Currently working on Heptio Ark backup software

# Administrative Stuff

- Format: two main parts of two subparts each (concepts and practice).
  - We'll take couple of human-needs breaks
  - If you have already been playing with either of the technologies here, you might not get much value out of this workshop
- 

# Our Q + A “rules”



- Do you know the answer to this question?
- Is it really a comment but with a question mark tacked on?
- Is the question rhetorical?
- Trying to show off how much you know about the subject?

# Our Q + A “rules”



- Speed of delivery
- Not understanding
- Request to repeat explanation
- Anything immediately related to the topic(s) at hand

# DON'T PANIC




Why this  
Workshop?



Container adoption has been insane in the last 4 years.



## Containerization timeline




SOURCE: GARTNER


However, containers by themselves are an incomplete story...

# Containers


Need a story for ...

- Deployment at scale
- Orchestration
- Monitoring
- Efficient host resources usage






- Tightly coupled to Docker
- Too opinionated
- Difficult to extend
- Internal networking architecture hard to interop
- Project open source but governed by Docker



- Not container specific
- Substantial learning curve
- Scheduler too generic
- Major parts written in different languages
- Overkill for smaller deployments



**kubernetes**







**kubernetes**

- Solved a lot of the issues other orchestrators had
- Gentler learning curve
- Container runtime independent
- ... but much more detail about this later

# Part I. Containers

Containers are  
not really a new  
thing.

# Abridged History

2000 “Jails” are included in FreeBSD.

2001-ish, the precursor to Virtual Private Server is introduced to the Linux ecosystem

2004 Sun adds “Zones” to Solaris.

2006-ish early precursors of cgroups raise to the fore

2008 Kernel namespaces and LXC (Linux Containers)

2013 Docker is released

# Runtime Alternatives



CRI-O: OCI-based  
Kubernetes Runtime



# Main Benefits

- Velocity
- Portability\*
- Reliability
- Efficiency
- Isolation
- Infrastructure as code

\*mostly

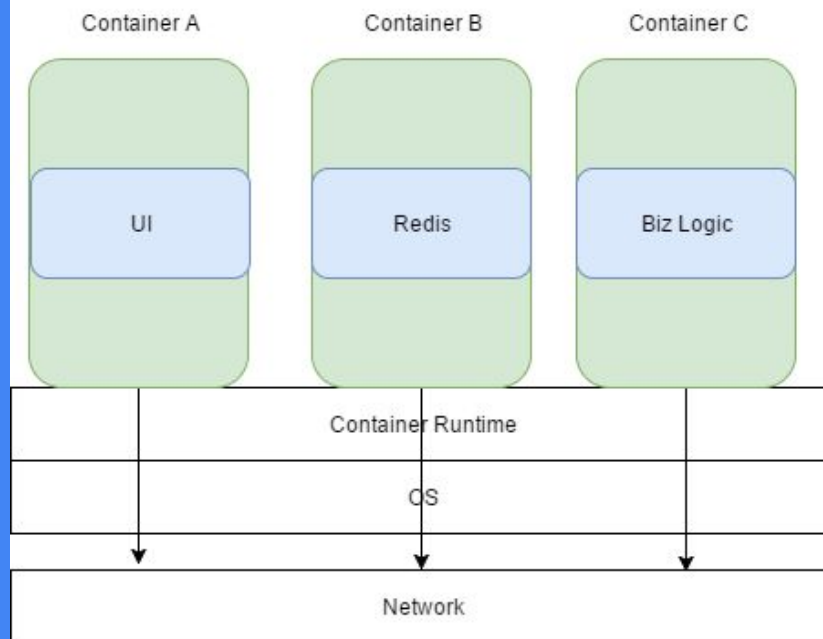
# Containers aren't VMs

- Full OS on virtualized hardware
- “Hypervisor Tax”
- Provisioning can take a long time (relatively speaking)
- Images are immutable and monolithic.
- Substantial config work
- Decouples application from underlying OS
- Composable Images
- Lean
- Very fast provisioning times
- Very “Disposable”



“All great stuff,  
but now with less  
hand waving, pls”

# Bird's Eye View



# The Container Model

- Shares resources with other processes, but constrained to high degree of isolation
- Can access persistent storage as volumes
- Can share/pass context

# Container Images


- Vaguely similar in concept with VM images, OS images, etc.
- In Containerland an image is a bundle of layers that together form only logical unit
- 1:N (image:containers)

# Container Images

- In the Docker ecosystem they are named and labeled as follows:  
`<user>/<imagename>:<tag>`  
Ex: ruben/introtock:latest (not an actual thing)
- Generally speaking images default to 'latest'

# Container Images

- Almost every instruction in a Dockerfile creates a layer in the image



```
FROM php:5-apache
```

```
RUN apt-get update
```

```
RUN apt-get install -y  
php-pear
```

```
RUN pear channel-discover
```

# Container Images

- The layered file system is key to images being composable
  - Key difference between container and traditional VM/OS images
  - All layers are presented as one logical unit


# Image Registries

- Registries are network storage for images
- You can publish (push) to or download (pull) from them
- Registries can be public, private and local



# Image Registries

``docker run user/image:tag``



Let's hit the shell

# Quiz Time!

- Create a new directory
- Create Dockerfile that builds on python:2 and prints to stdout The Zen of Python (hint 'this' module)
  - Hint: it only takes 2 lines (one for FROM and one for CMD)
- Build/tag the image
- Push Image to local registry
- Delete image from local cache
- Run container

“I thought you  
said Kubernetes.  
Where’s the  
Kubernetes?”

# The Motivation

- Google calculated “human errors” were root cause of over 99% of outages
- Continuous system changes
- Lack of deployment consistency and scalability

- Google creates Borg to address these issues for their own infrastructure
- Insane learning curve
- Engineer and operator ramp-up was (still is!) lengthy

- But 99.99% of users aren't Google
  - infinitesimal scale in comparison
  - Budget, talent, etc.
  - Line of business



- A group of Googlers (Joe\*, Brendan, et al) saw the opportunity
  - General applicability
  - Open Source

Full disclosure: \* my employer

# Design Principles

- Strong focus on developer and operator experience
- Easily extensible, with few opinions
- A subset of Borg, but for the masses

# What can it do?

- Aggregate virtual and bare metal infrastructure into a cluster
- Container orchestration within clusters
- Provide straightforward app lifecycle management

# Key Features


- Automatic binpacking
- Straightforward horizontal scaling
- Automatable rollouts/rollbacks
- Deployment “medic”
- Highly Fault Tolerant (including itself!)

# Key Features

- Service discovery
- Load balancing
- Secret management
- Batch execution

“Nice. But would  
like more  
specifics, though”


## Orbital View



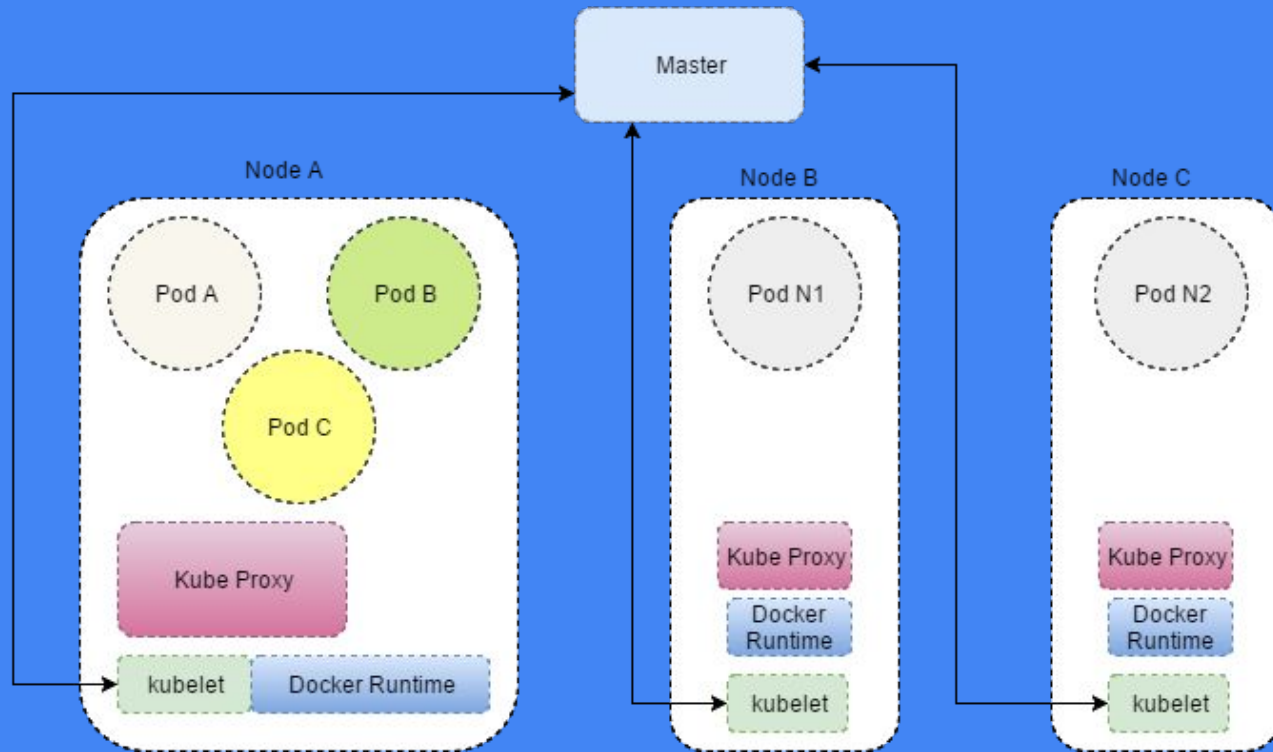
Zoom...enhance....




# Master Node



## Worker Nodes "Node" Zoom



# Networking



# Networking

- Scenario A: pod-to-pod (flat network, basic IP)
- Scenario B: Pod-to-Service (virtual-IP, through kubeproxy)
- Scenario C: External-to-service (fronted by a LB -- not shown -- and routed by kubeproxy)

# Main Components: Objects

- Objects are persistent entities that manage and represent the state of certain parts of the system
  - Node
  - Service
  - Pod
  - Deployment ... etc.


# Main Components: Objects

- Objects describe state of:
  - What deployments/services are running
  - Node status
  - Services, storage
  - Deployments
  - Namespaces ... etc.

# Pods

- Smallest unit of compute that can be managed by Kubernetes
  - Containers within pods run as if they're in a single host
  - Share namespaces, IP addr and port space
  - Comms through localhost:port or over IPC
  - Mount/use same volumes

# Labels

- Key/value pairs that are associated to objects and carry semantic value within Kubernetes
  - Organize and select resources
  - Commonly used to ID releases, environment, tiers
  - Can be used as query selectors
- 



# Annotations

- Somewhat similar to Labels, except:
  - Meaningless to Kubernetes
  - Can't be queried (i.e. not indexed)
  - Usually good for bespoke/extended metadata

# Services

- A grouping of Pods running in on the cluster
  - Load balancing
  - Service Discovery
  - Zero downtime deployments/upgrades

# ReplicaSets

- Ensures specified number of pod instances are running at any given time.
- Can update many aspects via spec definition
- Can manage one or more Pods based on labels

# Deployments

- Manages ReplicaSets
- Best object for deploying and managing applications in Kubernetes
- Manages the update process
- More

Back to the shell ...

Before we go ...



# Don't Forget Security

- Secrets
  - Hold sensitive data (credentials, keys, etc.)
  - RBAC/ABAC configurable
  - Can be leveraged by the system or regular users

# Don't Forget Security

- Authorization:
  - RBAC: system auth managed by Kubernetes
    - Impose restrictions at the Namespace level
  - WARNING: comes “disabled” by default -- not production ready in some bootstrappers.



# The End.

Thank you!

Tweet at us @rdodev @palendae